



МИНОБРНАУКИ РОССИИ
Федеральное государственное бюджетное образовательное учреждение высшего
образования

«МИРЭА – Российский технологический университет»

РТУ МИРЭА

ИКБ направление «Киберразведка и противодействие угрозам с
применением технологий искусственного интеллекта» 10.04.01

Кафедра КБ-4 «Интеллектуальные системы информационной
безопасности»

Лабораторная работа №1

по дисциплине

«Анализ защищенности систем искусственного интеллекта»

Группа:
ББМО-01-22

Выполнил:
Дмитриев М.Н.

Проверил:
Спирин А.А.

Москва 2023

- Клонирование репозитория.

```
In [1]: !git clone https://github.com/ewatson2/EEL6812_DeepFool_Project.git
```

```
Cloning into 'EEL6812_DeepFool_Project'...
remote: Enumerating objects: 96, done.
remote: Counting objects: 100% (3/3), done.
remote: Compressing objects: 100% (2/2), done.
remote: Total 96 (delta 2), reused 1 (delta 1), pack-reused 93
Receiving objects: 100% (96/96), 33.99 MiB | 15.11 MiB/s, done.
Resolving deltas: 100% (27/27), done.
```

- Сменим директорию

```
%cd /content/EEL6812_DeepFool_Project
```

```
/content/EEL6812_DeepFool_Project
```

- Импортируем библиотеки.

```
import numpy as np
import os
import json, torch
from torch.utils.data import DataLoader, random_split
from torchvision import datasets, models
from torchvision.transforms import transforms
from models.project_models import FC_500_150, LeNet_CIFAR, LeNet_MNIST, Net
from utils.project_utils import get_clip_bounds, evaluate_attack, display_attack
```

- Установка случайного значения, равного номеру в списке (39).

```
In [4]: rand_seed = 39
np.random.seed(rand_seed)
torch.manual_seed(rand_seed)

use_cuda = torch.cuda.is_available()
device = torch.device('cuda' if use_cuda else 'cpu')
```

- Загрузка датасета MNIST.

```
mnist_mean = 0.5
mnist_std = 0.5
mnist_dim = 28

mnist_min, mnist_max = get_clip_bounds(mnist_mean, mnist_std, mnist_dim)
mnist_min = mnist_min.to(device)
mnist_max = mnist_max.to(device)

mnist_tf = transforms.Compose([ transforms.ToTensor(), transforms.Normalize( mean=mnist_mean, std=mnist_std)])

mnist_tf_train = transforms.Compose([ transforms.RandomHorizontalFlip(), transforms.ToTensor(), transforms.Normalize( mean=mnist_mean, std=mnist_std)])

mnist_tf_inv = transforms.Compose([ transforms.Normalize( mean=0.0, std=np.divide(1.0, mnist_std)), transforms.Normalize( mean=np.multiply(-1.0, mnist_std), std=1.0)])

mnist_temp = datasets.MNIST(root='datasets/mnist', train=True, download=True, transform=mnist_tf_train)
mnist_train, mnist_val = random_split(mnist_temp, [50000, 10000])
mnist_test = datasets.MNIST(root='datasets/mnist', train=False, download=True, transform=mnist_tf)
```

- Загрузка датасета CIFAR-10.

```
cifar_mean = [0.491, 0.482, 0.447]
cifar_std = [0.202, 0.199, 0.201]
cifar_dim = 32

cifar_min, cifar_max = get_clip_bounds(cifar_mean, cifar_std, cifar_dim)
cifar_min = cifar_min.to(device)
cifar_max = cifar_max.to(device)
cifar_tf = transforms.Compose([ transforms.ToTensor(), transforms.Normalize( mean=cifar_mean, std=cifar_std)])

cifar_tf_train = transforms.Compose([ transforms.RandomCrop( size=cifar_dim, padding=4), transforms.RandomHorizontalFlip(), transforms.ToTensor(), transforms.Normalize( mean=cifar_mean, std=cifar_std)])

cifar_tf_inv = transforms.Compose([ transforms.Normalize( mean=[0.0, 0.0, 0.0], std=np.divide(1.0, cifar_std)), transforms.Normalize( mean=np.multiply(-1.0, cifar_mean), std=[1.0, 1.0, 1.0])])

cifar_temp = datasets.CIFAR10(root='datasets/cifar-10', train=True, download=True, transform=cifar_tf_train)

cifar_train, cifar_val = random_split(cifar_temp, [40000, 10000])

cifar_test = datasets.CIFAR10(root='datasets/cifar-10', train=False, download=True, transform=cifar_tf)

cifar_classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

- Выполнение настройки и загрузки DataLoader.

```
batch_size = 64
workers = 4
mnist_loader_train = DataLoader(mnist_train, batch_size=batch_size, shuffle=True, num_workers=workers)
mnist_loader_val = DataLoader(mnist_val, batch_size=batch_size, shuffle=False, num_workers=workers)
mnist_loader_test = DataLoader(mnist_test, batch_size=batch_size, shuffle=False, num_workers=workers)
cifar_loader_train = DataLoader(cifar_train, batch_size=batch_size, shuffle=True, num_workers=workers)
cifar_loader_val = DataLoader(cifar_val, batch_size=batch_size, shuffle=False, num_workers=workers)
cifar_loader_test = DataLoader(cifar_test, batch_size=batch_size, shuffle=False, num_workers=workers)
```

- Настройка обучающей модель.

```
train_model = True

epochs = 50
epochs_nin = 100

lr = 0.004
lr_nin = 0.01
lr_scale = 0.5

momentum = 0.9

print_step = 5

deep_batch_size = 10
deep_num_classes = 10
deep_overshoot = 0.02
deep_max_iters = 50

deep_args = [deep_batch_size, deep_num_classes, deep_overshoot, deep_max_iters]

if not os.path.isdir('weights/deepfool'): os.makedirs('weights/deepfool', exist_ok=True)

if not os.path.isdir('weights/fgsm'): os.makedirs('weights/fgsm', exist_ok=True)
```

- Загрузка и оценка стойкости модели Network-In-Network Model к FGSM и DeepFool атакам на основе датасета CIFAR-10.

```
fgsm_eps = 0.2
model = Net().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_nin.pth', map_location=torch.device('cpu')))
evaluate_attack('cifar_nin_fgsm.csv', 'results', device, model, cifar_loader_test, cifar_min, cifar_max, fgsm_eps, is_fgsm=True)
print('')
evaluate_attack('cifar_nin_deepfool.csv', 'results', device, model, cifar_loader_test, cifar_min, cifar_max, deep_args, is_fgsm=False)
if device.type == 'cuda': torch.cuda.empty_cache()

FGSM Test Error : 81.29%
FGSM Robustness : 1.77e-01
FGSM Time (All Images) : 0.67 s
FGSM Time (Per Image) : 67.07 us

DeepFool Test Error : 93.76%
DeepFool Robustness : 2.12e-02
DeepFool Time (All Images) : 185.12 s
DeepFool Time (Per Image) : 18.51 ms
```

- Загрузим и оценим стойкость модели LeNet к FGSM и DeepFool атакам на основе датасета CIFAR-10.

```
fgsm_eps = 0.1
model = LeNet_CIFAR().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_lenet.pth', map_location=torch.device('cpu')))
evaluate_attack('cifar_lenet_fgsm.csv', 'results', device, model, cifar_loader_test, cifar_min, cifar_max, fgsm_eps, is_fgsm=True)
print('')
evaluate_attack('cifar_lenet_deepfool.csv', 'results', device, model, cifar_loader_test, cifar_min, cifar_max, deep_args, is_fgsm=False)
if device.type == 'cuda': torch.cuda.empty_cache()

FGSM Test Error : 91.71%
FGSM Robustness : 8.90e-02
FGSM Time (All Images) : 0.40 s
FGSM Time (Per Image) : 40.08 us

DeepFool Test Error : 87.81%
DeepFool Robustness : 1.78e-02
DeepFool Time (All Images) : 73.27 s
DeepFool Time (Per Image) : 7.33 ms
```











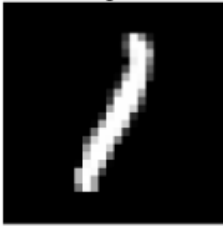

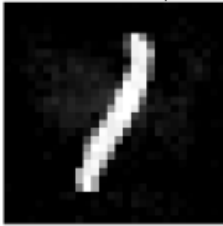

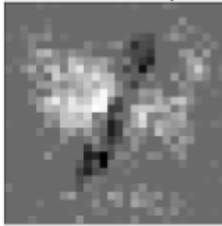




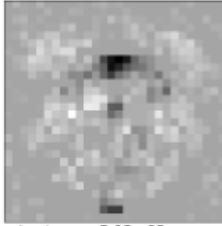
- Выполним оценку атакующих примеров для сетей.

```
#LeNet
fgsm_eps = 0.6
model = LeNet_MNIST().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_lenet.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_
if device.type == 'cuda': torch.cuda.empty_cache()


#FCNet
fgsm_eps = 0.2
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_
if device.type == 'cuda': torch.cuda.empty_cache()

#Network-in-Network
fgsm_eps = 0.2
model = Net().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_nin.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_
if device.type == 'cuda': torch.cuda.empty_cache()

#LeNet CIFAR-10
fgsm_eps = 0.1
model = LeNet_CIFAR().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_lenet.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_
if device.type == 'cuda': torch.cuda.empty_cache()
```

Original	Adversarial (FGSM)	Adversarial (DeepFool)	Perturbation (FGSM)	Perturbation (DeepFool)
 <p>true label : 2 pred label : 2 conf score : 15.37</p>	 <p>pred label : 9 conf score : 15.30</p>	 <p>pred label : 3 conf score : 11.81</p>	 <p>robustness : 4.64e-01 eps : 0.6</p>	 <p>robustness : 5.01e-02 overshoot : 0.02 iters : 8</p>
 <p>true label : 2 pred label : 0 conf score : 9.73</p>	 <p>pred label : 0 conf score : 7.94</p>	 <p>pred label : 2 conf score : 9.51</p>	 <p>robustness : 4.41e-01 eps : 0.6</p>	 <p>robustness : 2.95e-03 overshoot : 0.02 iters : 6</p>
Original	Adversarial (FGSM)	Adversarial (DeepFool)	Perturbation (FGSM)	Perturbation (DeepFool)
 <p>true label : 1 pred label : 1 conf score : 15.27</p>	 <p>pred label : 4 conf score : 10.00</p>	 <p>pred label : 4 conf score : 10.84</p>	 <p>robustness : 1.45e-01 eps : 0.2</p>	 <p>robustness : 6.65e-02 overshoot : 0.02 iters : 9</p>
 <p>true label : 9 pred label : 9 conf score : 18.12</p>	 <p>pred label : 4 conf score : 14.10</p>	 <p>pred label : 4 conf score : 13.07</p>	 <p>robustness : 1.61e-01 eps : 0.2</p>	 <p>robustness : 5.93e-02 overshoot : 0.02 iters : 9</p>

\



- Отражаем отличия для fgsm_eps=(0.001, 0.02, 0.5, 0.9, 10).




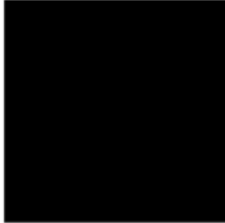
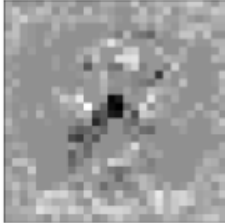




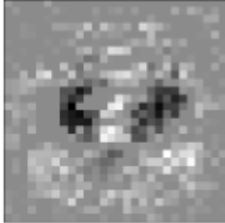




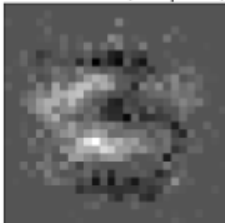




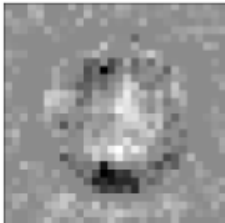
```
fgsm_eps = 0.001
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_
if device.type == 'cuda': torch.cuda.empty_cache()


fgsm_eps = 0.002
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 0.5
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 0.9
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 10
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_
if device.type == 'cuda': torch.cuda.empty_cache()
```

Original	Adversarial (FGSM)	Adversarial (DeepFool)	Perturbation (FGSM)	Perturbation (DeepFool)
 true label : 8 pred label : 8 conf score : 12.86	 pred label : 8 conf score : 12.79	 pred label : 9 conf score : 9.74	 robustness : 8.67e-04 eps : 0.001	 robustness : 4.25e-02 overshoot : 0.02 iters : 12
 true label : 9 pred label : 9 conf score : 20.62	 pred label : 9 conf score : 20.59	 pred label : 3 conf score : 13.74	 robustness : 8.58e-04 eps : 0.001	 robustness : 9.63e-02 overshoot : 0.02 iters : 10
 true label : 3 pred label : 3 conf score : 16.43	 pred label : 3 conf score : 16.32	 pred label : 8 conf score : 9.18	 robustness : 1.60e-03 eps : 0.002	 robustness : 6.94e-02 overshoot : 0.02 iters : 7
 true label : 0 pred label : 0 conf score : 20.63	 pred label : 0 conf score : 20.55	 pred label : 9 conf score : 12.24	 robustness : 1.71e-03 eps : 0.002	 robustness : 1.51e-01 overshoot : 0.02 iters : 10



- Проверим влияние параметра fgsm_eps для FC на датасете MNIST.



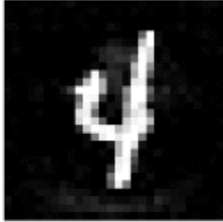

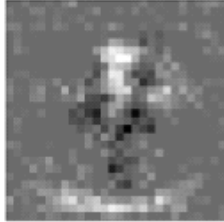




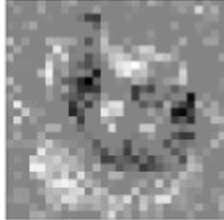

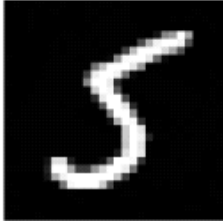
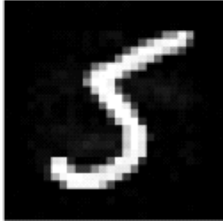

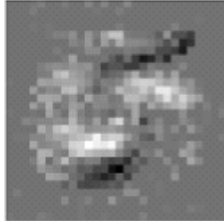

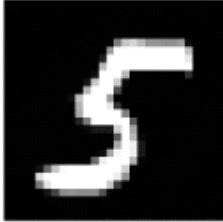
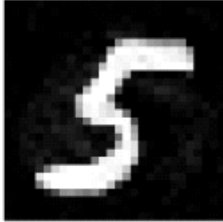

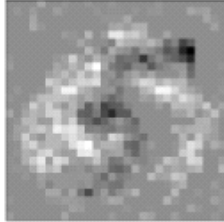
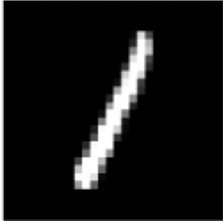

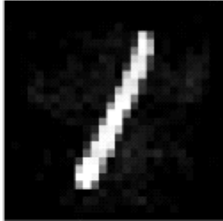

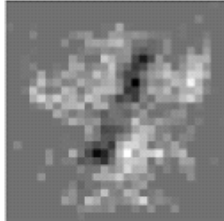
```
fgsm_eps = 0.001
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_labels=False, :
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 0.02
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_labels=False, :
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 0.5
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_labels=False, :
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 0.9
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_labels=False, :
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 10
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_labels=False, :
if device.type == 'cuda': torch.cuda.empty_cache()
```

Original	Adversarial (FGSM)	Adversarial (DeepFool)	Perturbation (FGSM)	Perturbation (DeepFool)
 true label : 4 pred label : 4 conf score : 15.39	 pred label : 4 conf score : 15.35	 pred label : 9 conf score : 10.13	 robustness : 8.27e-04 eps : 0.001	 robustness : 8.53e-02 overshoot : 0.02 iters : 9
 true label : 6 pred label : 6 conf score : 19.03	 pred label : 6 conf score : 18.97	 pred label : 2 conf score : 9.49	 robustness : 7.79e-04 eps : 0.001	 robustness : 1.13e-01 overshoot : 0.02 iters : 12
 true label : 5 pred label : 5 conf score : 13.76	 pred label : 5 conf score : 12.40	 pred label : 8 conf score : 7.80	 robustness : 1.52e-02 eps : 0.02	 robustness : 5.22e-02 overshoot : 0.02 iters : 8
 true label : 5 pred label : 5 conf score : 20.37	 pred label : 5 conf score : 18.91	 pred label : 2 conf score : 10.55	 robustness : 1.61e-02 eps : 0.02	 robustness : 8.17e-02 overshoot : 0.02 iters : 11
 true label : 1 pred label : 1 conf score : 15.09	 pred label : 4 conf score : 6.44	 pred label : 8 conf score : 7.99	 robustness : 3.59e-01 eps : 0.5	 robustness : 7.63e-02 overshoot : 0.02 iters : 10



true label : 4
pred label : 4
conf score : 18.09



pred label : 7
conf score : 20.71



pred label : 9
conf score : 10.21



robustness : 4.08e-01
eps : 0.5



robustness : 1.05e-01
overshoot : 0.02
iters : 9

Original

Adversarial (FGSM)

Adversarial (DeepFool)

Perturbation (FGSM)

Perturbation (DeepFool)



true label : 9
pred label : 9
conf score : 13.50




pred label : 4
conf score : 20.66



pred label : 4
conf score : 11.36




robustness : 6.84e-01
eps : 0.9



robustness : 3.83e-02
overshoot : 0.02
iters : 10



true label : 5
pred label : 5
conf score : 16.44



pred label : 2
conf score : 14.07



pred label : 8
conf score : 8.73



robustness : 6.79e-01
eps : 0.9



robustness : 6.39e-02
overshoot : 0.02
iters : 8

Original

Adversarial (FGSM)

Adversarial (DeepFool)

Perturbation (FGSM)

Perturbation (DeepFool)



true label : 3
pred label : 3
conf score : 22.63




pred label : 5
conf score : 61.76



pred label : 5
conf score : 14.52



robustness : 1.43e+00
eps : 10



robustness : 9.35e-02
overshoot : 0.02
iters : 10



true label : 0
pred label : 0
conf score : 11.48



pred label : 7
conf score : 56.61



pred label : 7
conf score : 8.74



robustness : 1.56e+00
eps : 10



robustness : 4.94e-02
overshoot : 0.02
iters : 10

- Проверим влияние параметра fgsm_eps для LeNet на датасете MNIST.





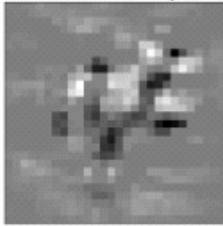




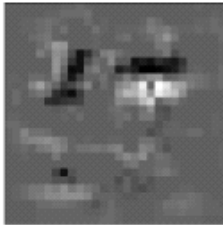
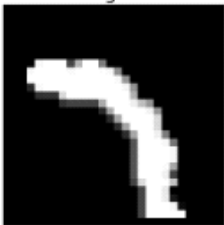

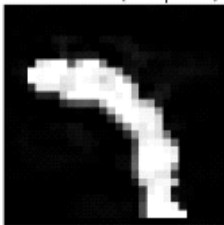

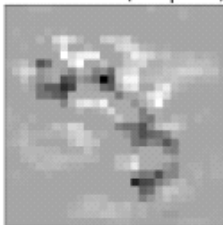




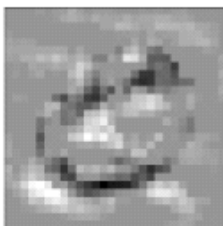
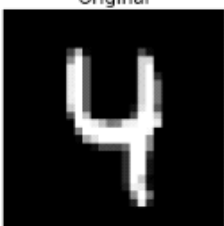

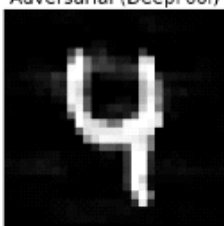

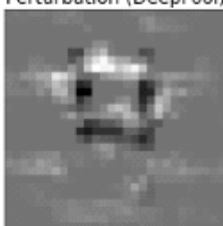
```
fgsm_eps = 0.001
model = LeNet_MNIST().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_lenet.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0.01)
if device.type == 'cuda': torch.cuda.empty_cache()


fgsm_eps = 0.02
model = LeNet_MNIST().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_lenet.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0.01)
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 0.5
model = LeNet_MNIST().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_lenet.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0.01)
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 0.9
model = LeNet_MNIST().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_lenet.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0.01)
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 10
model = LeNet_MNIST().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_lenet.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0.01)
if device.type == 'cuda': torch.cuda.empty_cache()
```


Original	Adversarial (FGSM)	Adversarial (DeepFool)	Perturbation (FGSM)	Perturbation (DeepFool)
 true label : 4 pred label : 4 conf score : 29.95	 pred label : 4 conf score : 29.95	 pred label : 9 conf score : 19.23	 robustness : 8.12e-04 eps : 0.001	 robustness : 1.24e-01 overshoot : 0.02 iters : 10
 true label : 5 pred label : 5 conf score : 12.22	 pred label : 5 conf score : 12.15	 pred label : 3 conf score : 8.95	 robustness : 8.28e-04 eps : 0.001	 robustness : 3.36e-02 overshoot : 0.02 iters : 9
Original	Adversarial (FGSM)	Adversarial (DeepFool)	Perturbation (FGSM)	Perturbation (DeepFool)
 true label : 7 pred label : 7 conf score : 15.81	 pred label : 7 conf score : 15.02	 pred label : 9 conf score : 11.57	 robustness : 1.57e-02 eps : 0.02	 robustness : 6.26e-02 overshoot : 0.02 iters : 10
 true label : 0 pred label : 0 conf score : 20.85	 pred label : 0 conf score : 20.36	 pred label : 9 conf score : 11.75	 robustness : 1.67e-02 eps : 0.02	 robustness : 1.91e-01 overshoot : 0.02 iters : 8
Original	Adversarial (FGSM)	Adversarial (DeepFool)	Perturbation (FGSM)	Perturbation (DeepFool)
 true label : 4 pred label : 4 conf score : 28.11	 pred label : 9 conf score : 26.76	 pred label : 9 conf score : 18.26	 robustness : 3.80e-01 eps : 0.5	 robustness : 8.73e-02 overshoot : 0.02 iters : 7



- Проверим влияние параметра fgsm_eps для NiN на датасете Cifar-10.

```
fgsm_eps = 0.001
model = Net().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_nin.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 0.02
model = Net().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_nin.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 0.5
model = Net().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_nin.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 0.9
model = Net().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_nin.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 10
model = Net().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_nin.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()
```




true label : truck
pred label : truck
conf score : 31.11




pred label : truck
conf score : 30.98



pred label : automobile
conf score : 28.35



robustness : $8.71e-04$
eps : 0.001



robustness : $1.41e-02$
overshoot : 0.02
iters : 2




true label : bird
pred label : bird
conf score : 16.76



pred label : bird
conf score : 16.66



pred label : cat
conf score : 11.04



robustness : $4.81e-04$
eps : 0.001



robustness : $1.86e-02$
overshoot : 0.02
iters : 4

Original

Adversarial (FGSM)

Adversarial (DeepFool)

Perturbation (FGSM)

Perturbation (DeepFool)




true label : cat
pred label : cat
conf score : 28.27




pred label : cat
conf score : 23.79



pred label : dog
conf score : 23.59



robustness : $1.83e-02$
eps : 0.02



robustness : $1.37e-02$
overshoot : 0.02
iters : 3



true label : horse
pred label : horse
conf score : 45.79



pred label : horse
conf score : 41.30



pred label : deer
conf score : 35.28



robustness : $1.70e-02$
eps : 0.02



robustness : $3.21e-02$
overshoot : 0.02
iters : 2

Original

Adversarial (FGSM)

Adversarial (DeepFool)

Perturbation (FGSM)

Perturbation (DeepFool)



true label : frog
pred label : frog
conf score : 34.21




pred label : bird
conf score : 18.70



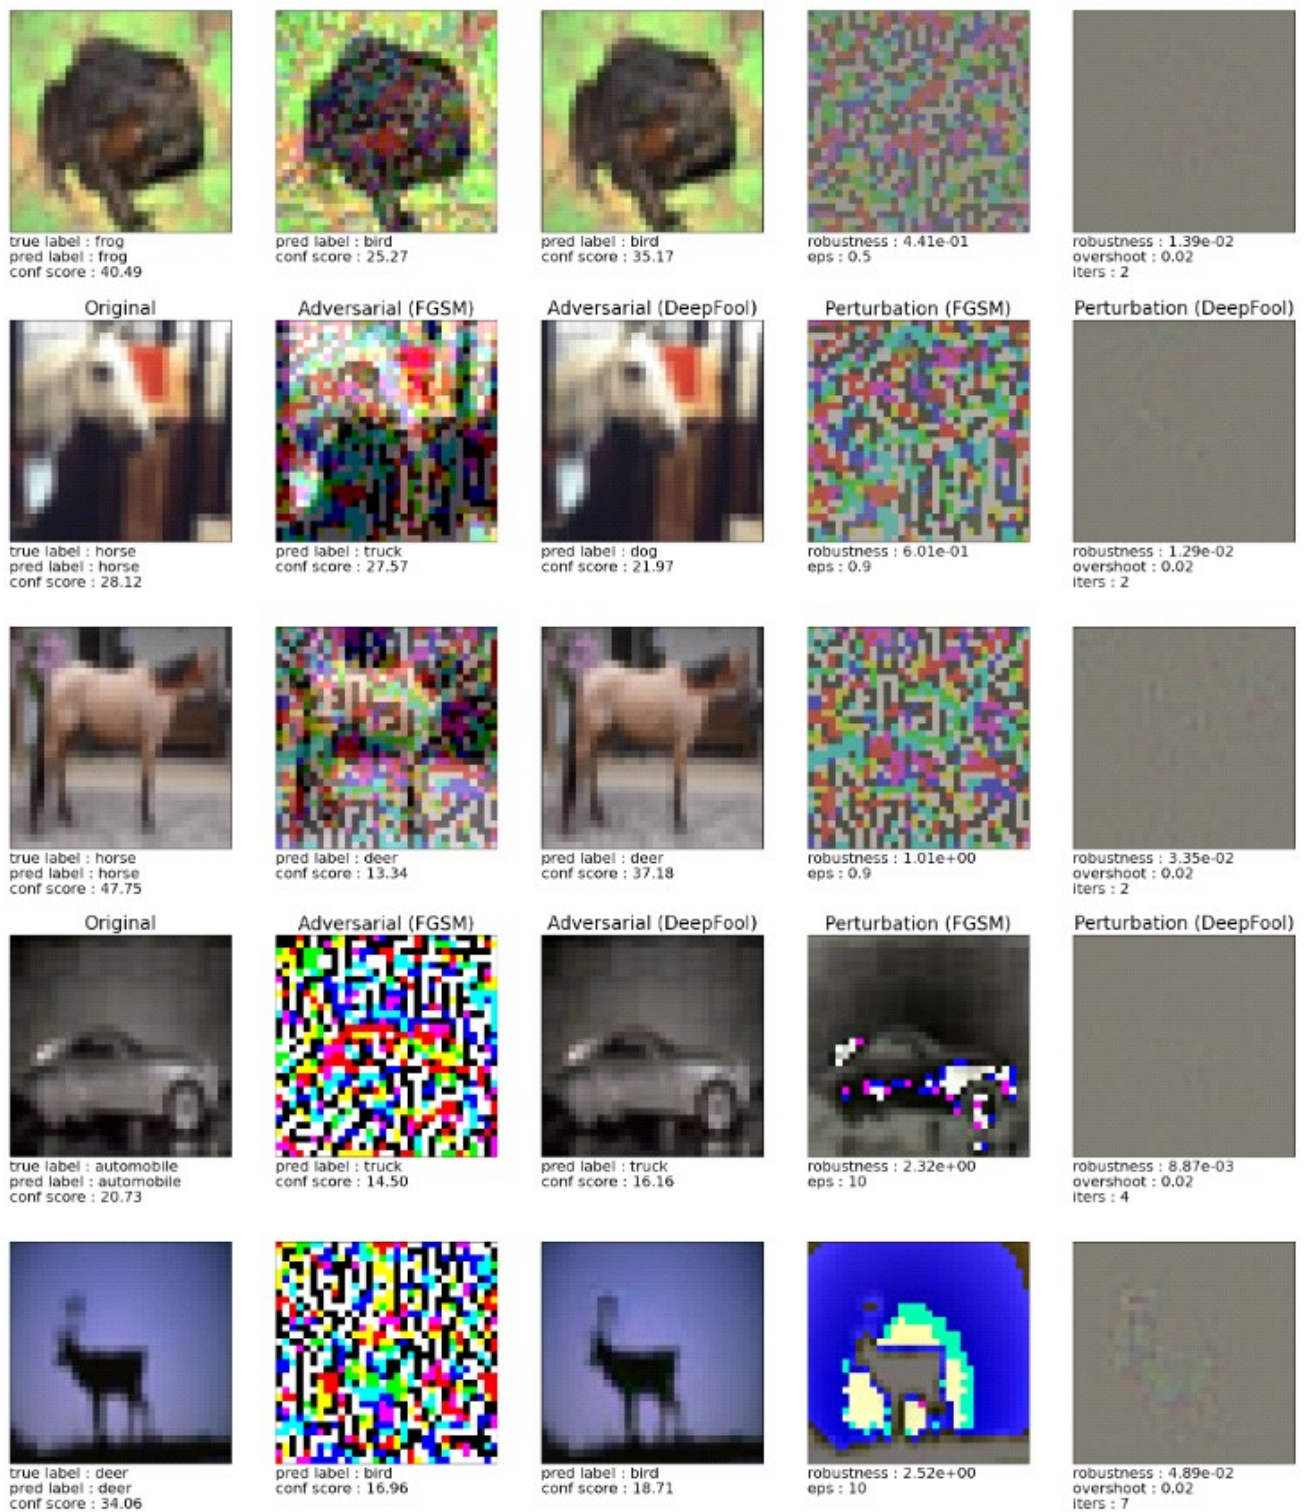
pred label : bird
conf score : 24.26



robustness : $7.10e-01$
eps : 0.5



robustness : $2.65e-02$
overshoot : 0.02
iters : 2



- Проверим влияние параметра fgsm_eps для LeNet на датасете Cifar-10.









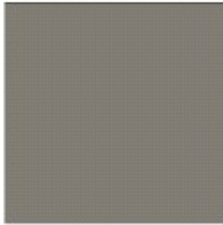
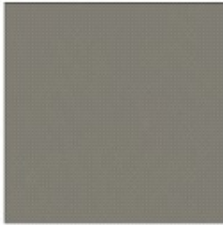









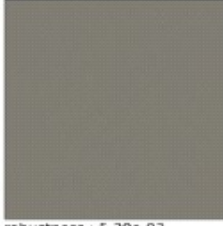



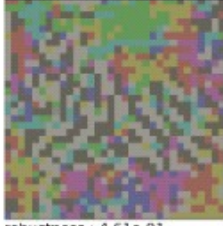


```
fgsm_eps = 0.001
model = LeNet_CIFAR().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_lenet.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

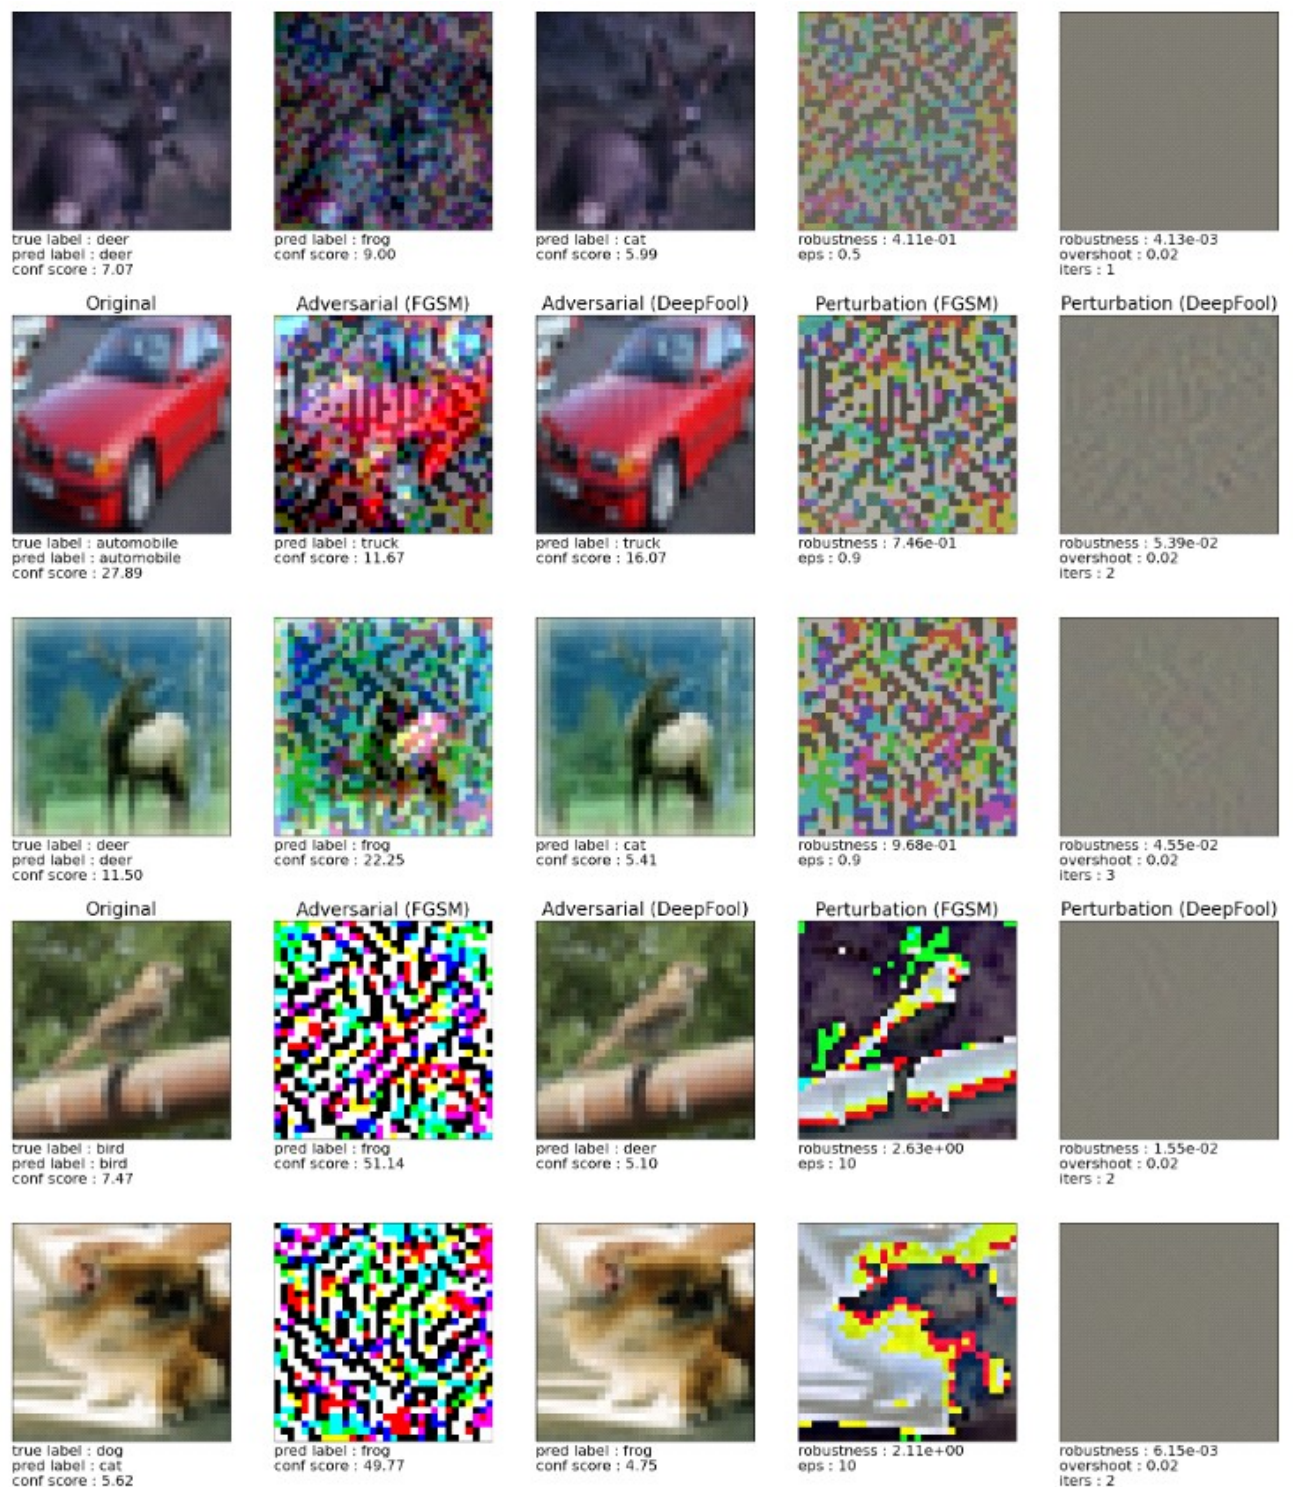
fgsm_eps = 0.02
model = LeNet_CIFAR().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_lenet.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 0.5
model = LeNet_CIFAR().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_lenet.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 0.9
model = LeNet_CIFAR().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_lenet.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 10
model = LeNet_CIFAR().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_lenet.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()
```

Original	Adversarial (FGSM)	Adversarial (DeepFool)	Perturbation (FGSM)	Perturbation (DeepFool)
 true label : deer pred label : deer conf score : 8.35	 pred label : deer conf score : 8.29	 pred label : horse conf score : 7.16	 robustness : $8.19\text{e-}04$ eps : 0.001	 robustness : $1.60\text{e-}02$ overshoot : 0.02 iters : 3
 true label : truck pred label : automobile conf score : 7.19	 pred label : automobile conf score : 7.32	 pred label : ship conf score : 6.05	 robustness : $7.32\text{e-}04$ eps : 0.001	 robustness : $4.48\text{e-}03$ overshoot : 0.02 iters : 2
Original	Adversarial (FGSM)	Adversarial (DeepFool)	Perturbation (FGSM)	Perturbation (DeepFool)
 true label : horse pred label : horse conf score : 24.45	 pred label : horse conf score : 21.08	 pred label : dog conf score : 13.60	 robustness : $1.83\text{e-}02$ eps : 0.02	 robustness : $4.95\text{e-}02$ overshoot : 0.02 iters : 3
 true label : truck pred label : cat conf score : 2.91	 pred label : cat conf score : 3.28	 pred label : dog conf score : 2.61	 robustness : $1.99\text{e-}02$ eps : 0.02	 robustness : $5.30\text{e-}03$ overshoot : 0.02 iters : 3
Original	Adversarial (FGSM)	Adversarial (DeepFool)	Perturbation (FGSM)	Perturbation (DeepFool)
 true label : ship pred label : ship conf score : 17.76	 pred label : bird conf score : 4.53	 pred label : airplane conf score : 10.46	 robustness : $4.61\text{e-}01$ eps : 0.5	 robustness : $3.11\text{e-}02$ overshoot : 0.02 iters : 2



Вывод: параметр fgsm_eps влияет на устойчивость сети. При увеличении значения fgsm_eps сети становятся уязвимыми к атакам и увеличивается ошибка тестирования модели. Это указывает на снижение ее производительности